

I'm Not ~~Dead~~ **undef** Yet

Perl: persistent subroutine memory allocation
(follow up to Cees' Lightning Talk)

Source code available at

https://github.com/dnstandish/TPM_perl_persistent_subroutine_memory_allocation

This presentation and actual output available
under releases

This exploration was primarily conducted with perl
5.18.2 on Linux

At the end I'll compare to perl 5.28.1

Persistence of memory allocation in functions

- storage allocated to lexical variable persists between calls

```
use Devel::Peek;  
$Devel::Peek::pv_limit = 16;  
sub big_string {  
    my $x;  
    Dump $x;  
    $x = "a" x 80000;  
    Dump $x;  
}
```

```
big_string;  
big_string;
```

SV = NULL(0x0) at 0x8edd3f8
REFCNT = 1
FLAGS = (PADMY)
SV = PV(0x8ec36b8) at 0x8edd3f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x8f48790 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004
SV = PV(0x8ec36b8) at 0x8edd3f8
REFCNT = 1
FLAGS = (PADMY)
PV = 0x8f48790 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004
SV = PV(0x8ec36b8) at 0x8edd3f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x8f48790 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004

Returning a value does not deallocate

```
sub big_string {  
  my $x;  
  Dump $x;  
  $x = "a" x 80000;  
  Dump $x;  
  return $x;  
}
```

```
my $x1 = big_string;  
Dump $x1;  
my $x2 = big_string;  
Dump $x2;
```

SV = NULL(0x0) at 0x854b3f8
REFCNT = 1
FLAGS = (PADMY)
SV = PV(0x85316b8) at 0x854b3f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x85b5df8 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004
SV = PV(0x85317a8) at 0x854b4e8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x85c9680 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004
SV = PV(0x85316b8) at 0x854b3f8
REFCNT = 1
FLAGS = (PADMY)
PV = 0x85b5df8 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004
SV = PV(0x85316b8) at 0x854b3f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x85b5df8 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004
SV = PV(0x85317b0) at 0x8566e40
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x85dcf08 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004

Returning a reference frees storage from function

```
sub big_string {  
  my $x;  
  Dump $x;  
  $x = "a" x 80000;  
  Dump $x;  
  return \$x;  
}
```

```
my $x1 = big_string;  
Dump $x1;  
my $x2 = big_string;
```

...

SV = PV(0xa0506b8) at 0xa06a3f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0xa0d4e00 "aaaaaaaaaaaaaaaaa"...
CUR = 80000
LEN = 80004

SV = IV(0xa06a4e4) at 0xa06a4e8
REFCNT = 1
FLAGS = (PADMY,ROK)
RV = 0xa06a3f8
SV = PV(0xa0506b8) at 0xa06a3f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0xa0d4e00 "aaaaaaaaaaaaaaaaa"...
CUR = 80000
LEN = 80004

SV = NULL(0x0) at 0xa04fb44
REFCNT = 1
FLAGS = (PADMY)

SV = PV(0xa050768) at 0xa04fb44
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0xa0e8688 "aaaaaaaaaaaaaaaaa"...
CUR = 80000
LEN = 80004

What about assigning undef?

```
sub big_string {  
  my $x;  
  Dump $x;  
  $x = "a" x 80000;  
  Dump $x;  
  $x = undef;  
  Dump $x;  
  return;  
}
```

```
big_string;  
big_string;
```

SV = NULL(0x0) at 0x912e3f8
 REFCNT = 1
 FLAGS = (PADMY)
SV = PV(0x91146b8) at 0x912e3f8
 REFCNT = 1
 FLAGS = (PADMY,POK,pPOK)
 PV = 0x9198de8 "aaaaaaaaaaaaaaaaa"...\\0
 CUR = 80000
 LEN = 80004
SV = PV(0x91146b8) at 0x912e3f8
 REFCNT = 1
 FLAGS = (PADMY)
 PV = 0x9198de8 "aaaaaaaaaaaaaaaaa"...\\0
 CUR = 80000
 LEN = 80004
SV = PV(0x91146b8) at 0x912e3f8
 REFCNT = 1
 FLAGS = (PADMY)
 PV = 0x9198de8 "aaaaaaaaaaaaaaaaa"...\\0
 CUR = 80000
 LEN = 80004
SV = PV(0x91146b8) at 0x912e3f8
 REFCNT = 1
 FLAGS = (PADMY,POK,pPOK)
 PV = 0x9198de8 "aaaaaaaaaaaaaaaaa"...\\0
 CUR = 80000
 LEN = 80004
SV = PV(0x91146b8) at 0x912e3f8
 REFCNT = 1
 FLAGS = (PADMY)
 PV = 0x9198de8 "aaaaaaaaaaaaaaaaa"...\\0
 CUR = 80000
 LEN = 80004

undef call works!

```
sub big_string {  
  my $x;  
  Dump $x;  
  $x = "a" x 80000;  
  Dump $x;  
  undef $x;  
  Dump $x;  
}  
big_string;  
big_string;
```

SV = NULL(0x0) at 0x9c803f8
REFCNT = 1
FLAGS = (PADMY)
SV = PV(0x9c666b8) at 0x9c803f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x9ceade0 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004

SV = PV(0x9c666b8) at 0x9c803f8
REFCNT = 1
FLAGS = (PADMY)
PV = 0

SV = PV(0x9c666b8) at 0x9c803f8
REFCNT = 1
FLAGS = (PADMY)
PV = 0

SV = PV(0x9c666b8) at 0x9c803f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x9ceade0 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 100012

SV = PV(0x9c666b8) at 0x9c803f8
REFCNT = 1
FLAGS = (PADMY)
PV = 0

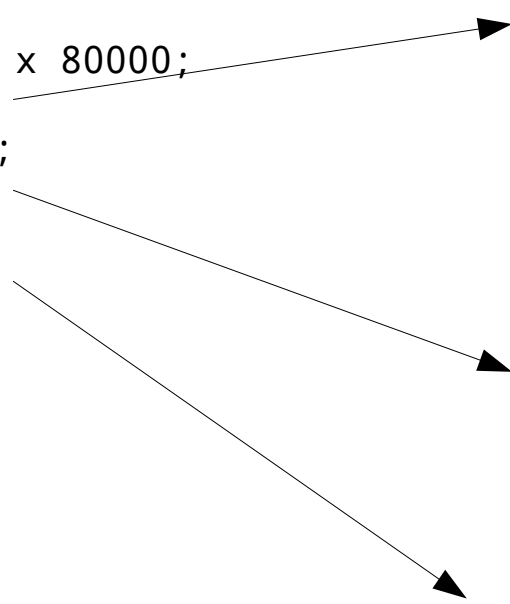
Assigning a value doesn't help

```
sub big_string {  
  my $x;  
  Dump $x;  
  $x = "a" x 80000;  
  Dump $x;  
  $x = "b";  
  Dump $x;  
  $x = 1;  
  Dump $x;  
}  
big_string;
```

SV = NULL(0x0) at 0x9a113f8
REFCNT = 1
FLAGS = (PADMY)
SV = PV(0x99f76b8) at 0x9a113f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x9a7bde8 "aaaaaaaaaaaaaaaaa"...\\0
CUR = 80000
LEN = 80004

SV = PV(0x99f76b8) at 0x9a113f8
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x9a7bde8 "b"\\0
CUR = 1
LEN = 80004

SV = PVIV(0x9a12ca0) at 0x9a113f8
REFCNT = 1
FLAGS = (PADMY,IOK,pIOK)
IV = 1
PV = 0x9a7bde8 "b"\\0
CUR = 1
LEN = 80004



local instead of lexical

```
use Devel::Peek;  
use feature 'say';  
$Devel::Peek::pv_limit = 16;
```

```
sub big_string {  
    my $offset = shift;  
    local $x;  
    Dump $x;  
    $x = (chr(ord("a") + $offset)) x (80000 - $offset  
* 16);  
    Dump $x;  
    return;  
}
```

```
big_string(1);
```

```
big_string(2);
```

SV = NULL(0x0) at 0x91948ec

REFCNT = 1

FLAGS = ()

SV = PV(0x91957a8) at 0x91948ec

REFCNT = 1

FLAGS = (POK,pPOK)

PV = 0x9216a88 "bbbbbbbbbbbbbbbbbb"... \0

CUR = 79984

LEN = 79988

SV = NULL(0x0) at 0x91948ec

REFCNT = 1

FLAGS = ()

SV = PV(0x91957a8) at 0x91948ec

REFCNT = 1

FLAGS = (POK,pPOK)

PV = 0x9216a88 "cccccccccccccccc"... \0

CUR = 79968

LEN = 79972

local seems to free up allocation between calls

Lots of functions with big variables will consume memory

- How to test?
- Devel::Peek::mstat()
 - not available
 - Devel::Peek::mstat: : perl not compiled with MYMALLOC
- on Linux can use /proc/<pid>/status

```
me@host:$ grep '^Vm' /proc/2471/status
```

```
VmPeak: 9156 kB  
VmSize: 9092 kB  
VmLck: 0 kB  
VmPin: 0 kB  
VmHWM: 4280 kB  
VmRSS: 4280 kB  
VmData: 2456 kB  
VmStk: 132 kB  
VmExe: 944 kB  
VmLib: 2160 kB  
VmPTE: 36 kB  
VmSwap: 0 kB
```

```
package ProcVM;

use autodie;

sub print_proc_vm {

    open my $fd, '<', "/proc/$$/status";
    while ( <$fd> ) {
        print if /^Vm/;
    }
    close $fd;

}

1;
```

use Template to generate script with multiple functions

```
...
[% FOREACH suffix IN suf_list %]
sub f[% suffix %] {
    my $offset = shift;
    print "sub [% suffix %]\n";
[% IF use_devel_peek %]
    Dump($s);
[% END %]
    my $s = (chr(ord("a") + [% suffix %] + $offset)) x
[% size %]
;
[% IF use_devel_peek %]
    Dump($s);
[% END %]
[% IF assign_undef %]
    $s = undef;
[% END %]
[% IF call_undef %]
    undef($s);
[% END %]
[% IF use_devel_peek %]
    Dump($s);
[% END %]
    return "";
}
[% END %]

[% FOREACH i IN call_list %]
print "cycle [% i %]\n";
ProcVM::print_proc_vm();
[% FOREACH suffix IN suf_list %]
f[% suffix %]([% i %]);
ProcVM::print_proc_vm();
sleep(1);
[% END %]
[% END %]
```

```

    ...
sub f1 {
    my $offset = shift;
    print "sub 1\n";
    my $s = (chr(ord("a") + 1 + $offset)) x 800000;
    return "";
}

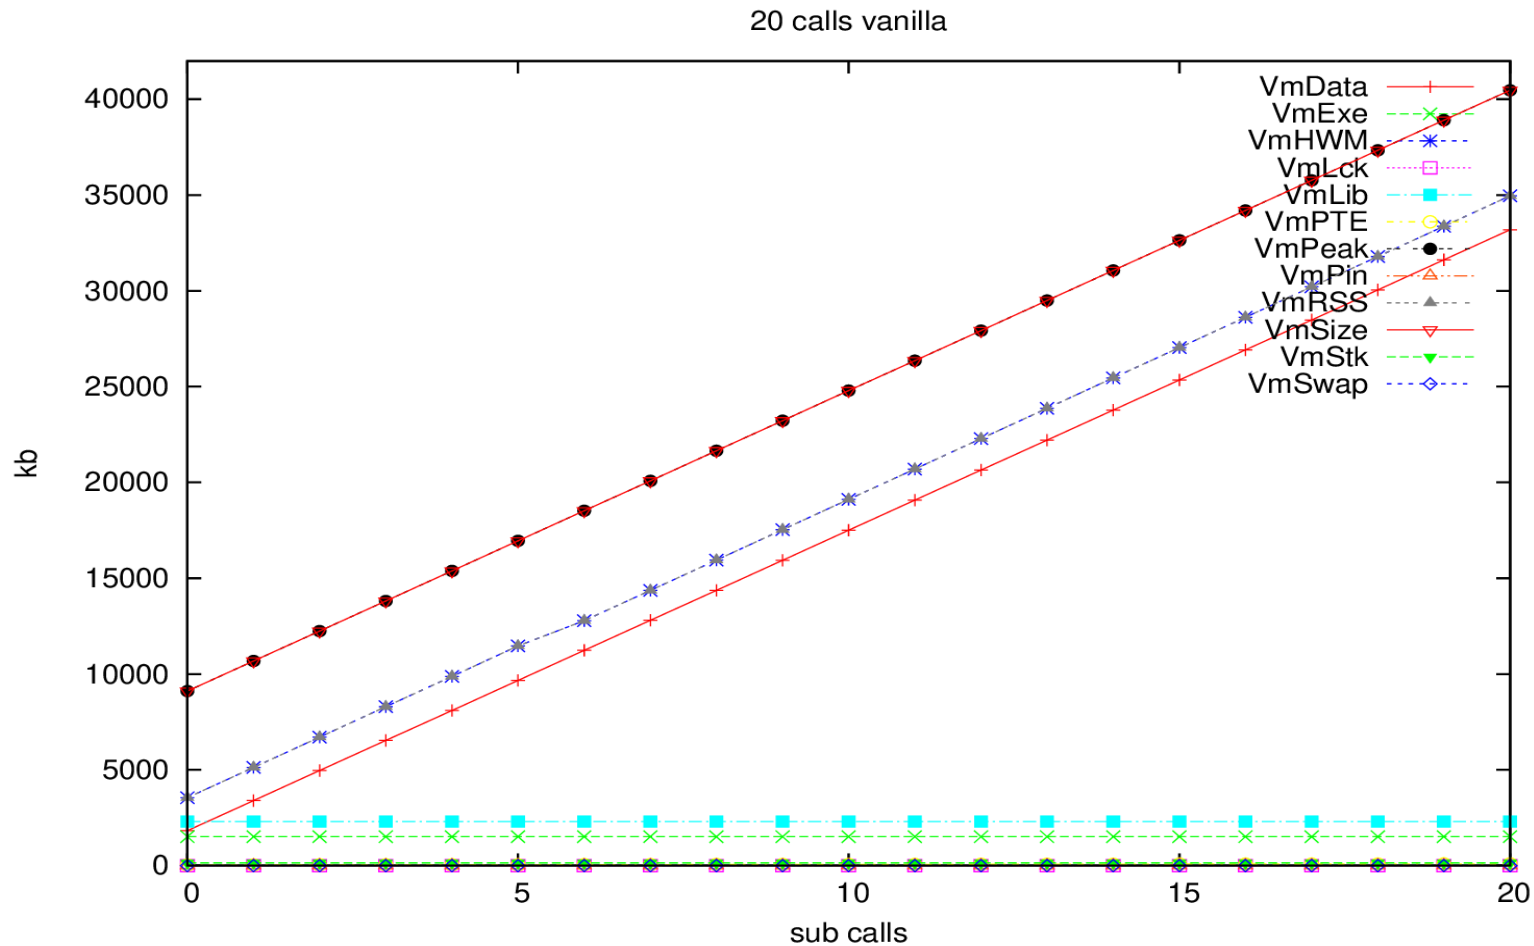
    ...
sub f20 {
    my $offset = shift;
    print "sub 20\n";
    my $s = (chr(ord("a") + 20 + $offset)) x 800000;
    return "";
}
print "cycle 1\n";
ProcVM::print_proc_vm();
f1(1);
ProcVM::print_proc_vm();

    ...
sleep(1);
f20(1);
ProcVM::print_proc_vm();
sleep(1);

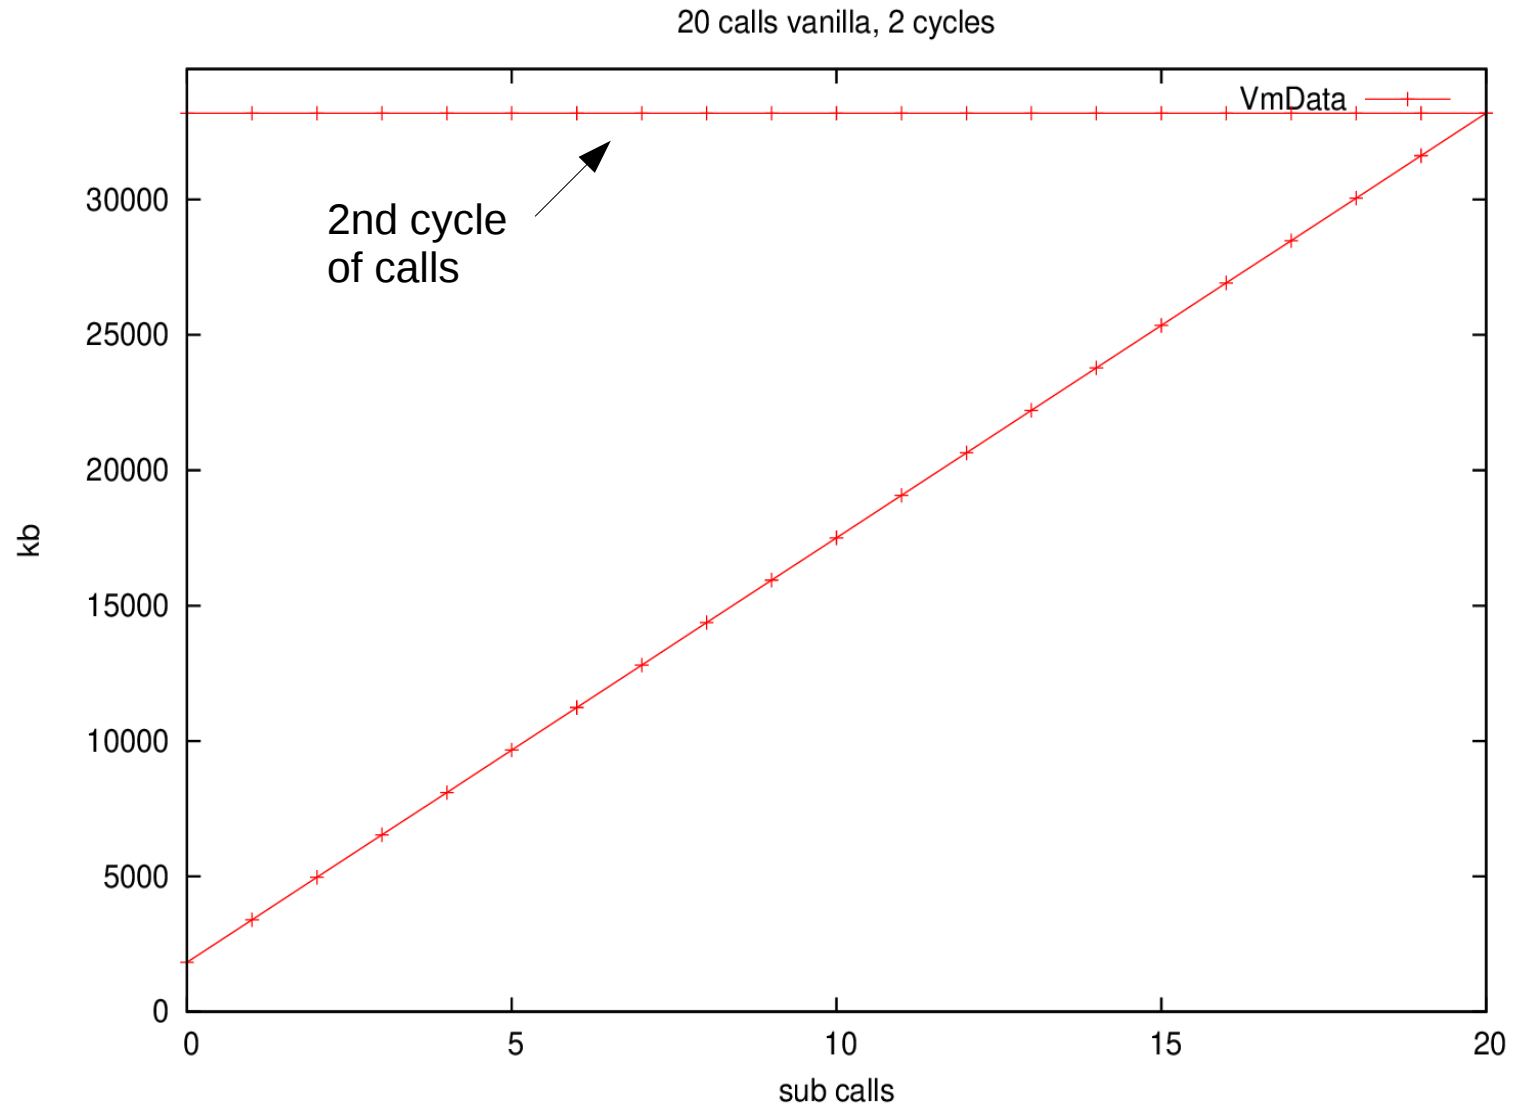
```

Parse output and use Chart::Gnuplot

per function memory consumption reflected in VmData, VmHWM, VMPeak, VmSize, VmRss



This is not a leak: calling function more than once does not consume additional memory



What happens if perl runs into a memory limit?

Does it reuse storage allocated to other functions? No, it dies.

./aa.out 16000000 func1.pl

setrlimit ok

execing func1.pl

template_file: many_func1.template

use_devel_peek:

size: 800000

n_func: 20

n_call: 1

assign_undef: 0

call_undef: 0

cycle 1

VmPeak: 9112 kB

VmSize: 9112 kB

VmLck: 0 kB

VmPin: 0 kB

VmHWM: 3528 kB

VmRSS: 3528 kB

VmData: 1828 kB

VmStk: 136 kB

VmExe: 1508 kB

VmLib: 2292 kB

VmPTE: 28 kB

VmSwap: 0 kB

sub 1

VmPeak: 10680 kB

VmSize: 10680 kB

VmLck: 0 kB

VmPin: 0 kB

VmHWM: 5112 kB

VmRSS: 5112 kB

VmData: 3396 kB

VmStk: 136 kB

VmExe: 1508 kB

VmLib: 2292 kB

VmPTE: 32 kB

VmSwap: 0 kB

.

.

.

sub 4

VmPeak: 15384 kB

VmSize: 15384 kB

VmLck: 0 kB

VmPin: 0 kB

VmHWM: 9864 kB

VmRSS: 9864 kB

VmData: 8100 kB

VmStk: 136 kB

VmExe: 1508 kB

VmLib: 2292 kB


VmPTE: 44 kB

VmSwap: 0 kB

sub 5

Out of memory!

Effect of calling undef on memory



Still see rise in memory use as each function called


Suspect memory is also allocated for temporary expression

```
sub f1 {  
  my $offset = shift;  
  print "sub 1\n";  
  my $s = (chr(ord("a") + 1 + $offset)) x 800000; ←  
  undef($s);  
  return "";  
}
```

What if expression is constant?

```
sub f1 {  
  my $offset = shift;  
  print "sub 1\n";  
  my $s = "a" x 800000; ←  
  undef($s);  
  return "";  
}
```

What if temporary is a constant? Is there optimization




- Yes, constant allocated before call
- not shared between functions

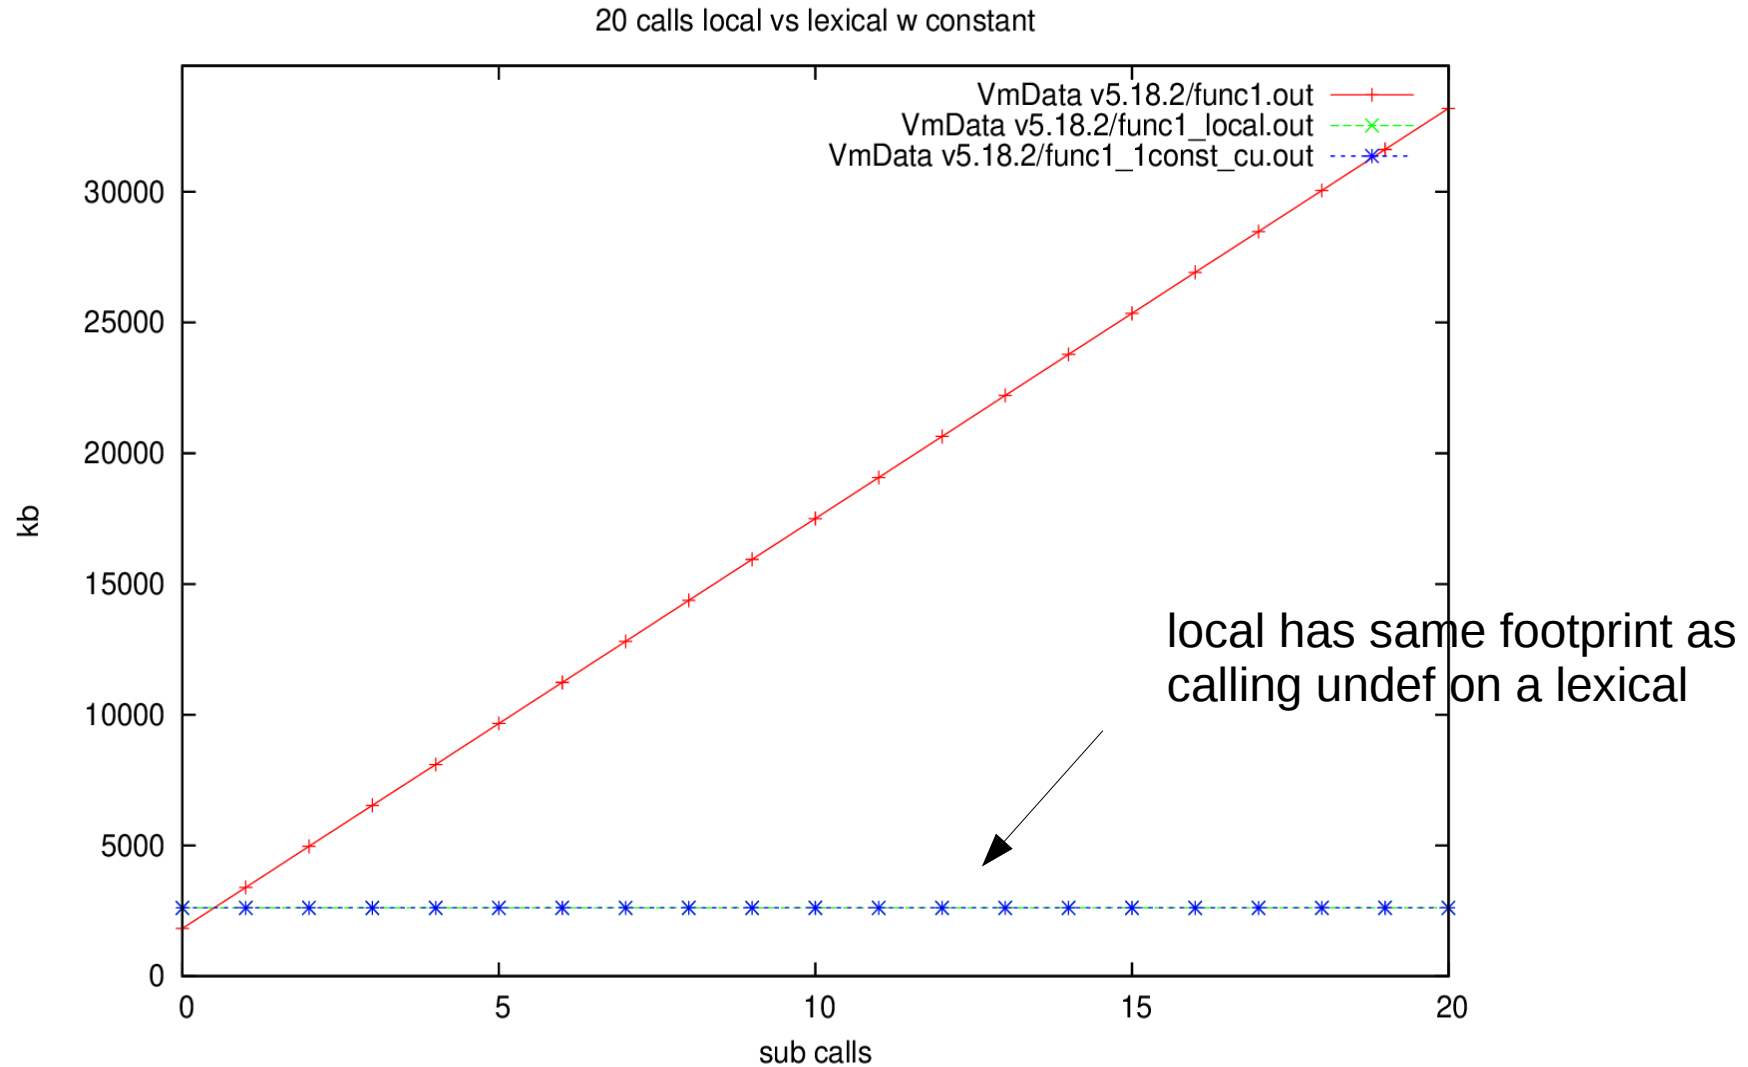
Put constant into a separate function

```
sub constant {  
    return "a" x 800000;  
}  
sub f1 {  
    my $offset = shift;  
    print "sub 1\n";  
    my $s = constant();  
  
    undef($s);  
    return "";  
}  
...
```

Segregating constant into function combined with calling undef minimizes memory use



Local with single constant



Closures

Perl closures shouldn't cross contaminate unintentionally

Suggests that closure use of sub lexical variable will prevent reuse of allocated storage

```

sub sub_gen {
    my $offset = shift;
    print STDERR "gen $offset\n";
    my $s;
    Dump $s;
    $s = (chr(ord("a") + $offset)) x 800000;
    Dump $s;
    return {
        info => sub {
            print STDERR "info $offset\n";
            Dump $s;
        },
        clear => sub {
            print STDERR "clear $offset\n";
            undef( $s );
        },
    }
}

my $c1 = sub_gen( 1 );
my $c2 = sub_gen( 2 );
$c1->{info}->();
$c1->{clear}->();
$c1->{info}->();
$c2->{info}->();
$c2->{clear}->();
$c2->{info}->();

```



```
my $c1 = sub_gen( 1 );
```

```
my $c2 = sub_gen( 2 );
```

gen 1

SV = NULL(0x0) at 0x8c015dc

REFCNT = 1

FLAGS = (PADMY)

SV = PV(0x8be77b0) at 0x8c015dc

REFCNT = 1

FLAGS = (PADMY,POK,pPOK)

PV = 0x406c2008 "bbbbbbbbbbbbbbbbbb"...\0

CUR = 800000

LEN = 800004

gen 2

SV = NULL(0x0) at 0x8c39e6c

REFCNT = 1

FLAGS = (PADMY)

SV = PV(0x8be77b8) at 0x8c39e6c

REFCNT = 1

FLAGS = (PADMY,POK,pPOK)

PV = 0x40786008 "cccccccccccccccc"...\0

CUR = 800000

LEN = 800004

Lexical variable \$s has different SV and PV for each call to sub_gen

gen 1
SV = NULL(0x0) at 0x8c015dc
REFCNT = 1
FLAGS = (PADMY)
SV = PV(0x8be77b0) at 0x8c015dc
REFCNT = 1
FLAGS = (PADMY,POK,pPOK)
PV = 0x406c2008 "bbbbbbbbbbbbbbbbbb"...\0
CUR = 800000
LEN = 800004

\$c1->{info}->();


info 1
SV = PV(0x8be77b0) at 0x8c015dc
REFCNT = 2
FLAGS = (PADMY,POK,pPOK)
PV = 0x406c2008 "bbbbbbbbbbbbbbbbbb"...\0
CUR = 800000
LEN = 800004

Closure takes ownership of lexical variable from generating function

Closure memory consumption test

- generate 10 closures
- generate 10 more, but invoke `clear->()` to undef closure variable

Closure memory use




Is Closure allocated memory released when ref count hits zero?

Expect allocated storage to be released once there are no more references

Otherwise this would be a source of memory leaks

- generate 10 closures and store in array
- empty the array
- generate 10 more closures

Closure memory is deallocated



Recursion

If functions hold on to allocated memory for lexical variables, how does this work with recursive functions?

```
sub get_const {  
    return chr(ord("a")) x 800000;  
}
```


```
sub recur {  
    my $cnt = shift;  
    my $limit = shift;  
    return if $cnt >= $limit;  
    $cnt++;  
    ProcVM::print_proc_vm();  
    print "sub $cnt\n";  
    my $s;  
    Dump $s;  
    $s = get_const();  
    Dump $s;  
    return recur($cnt, $limit);  
}
```

*opportunity for tail
recursion optimization*



```
print "cycle 1\n";  
recur(0, 20);  
recur(20, 45);  
ProcVM::print_proc_vm();
```


Memory allocated and retained independently by recursion depth




What if scalar isn't defined until after recursive call?

```
sub get_const {  
    return chr(ord("a")) x 800000;  
}
```

```
sub recur {  
    my $cnt = shift;  
    my $limit = shift;  
    return if $cnt >= $limit;  
    $cnt++;  
    ProcVM::print_proc_vm();  
    print "sub $cnt\n";  
    recur($cnt, $limit);  
    my $s;  
    Dump $s;  
    $s = get_const();  
    Dump $s;  
}
```

```
print "cycle 1\n";  
recur(0, 20);  
recur(20, 45);  
ProcVM::print_proc_vm();
```

Doesn't make much difference



Measurement of memory usage
precedes memory consumption.
Allocation not apparent until second
time through

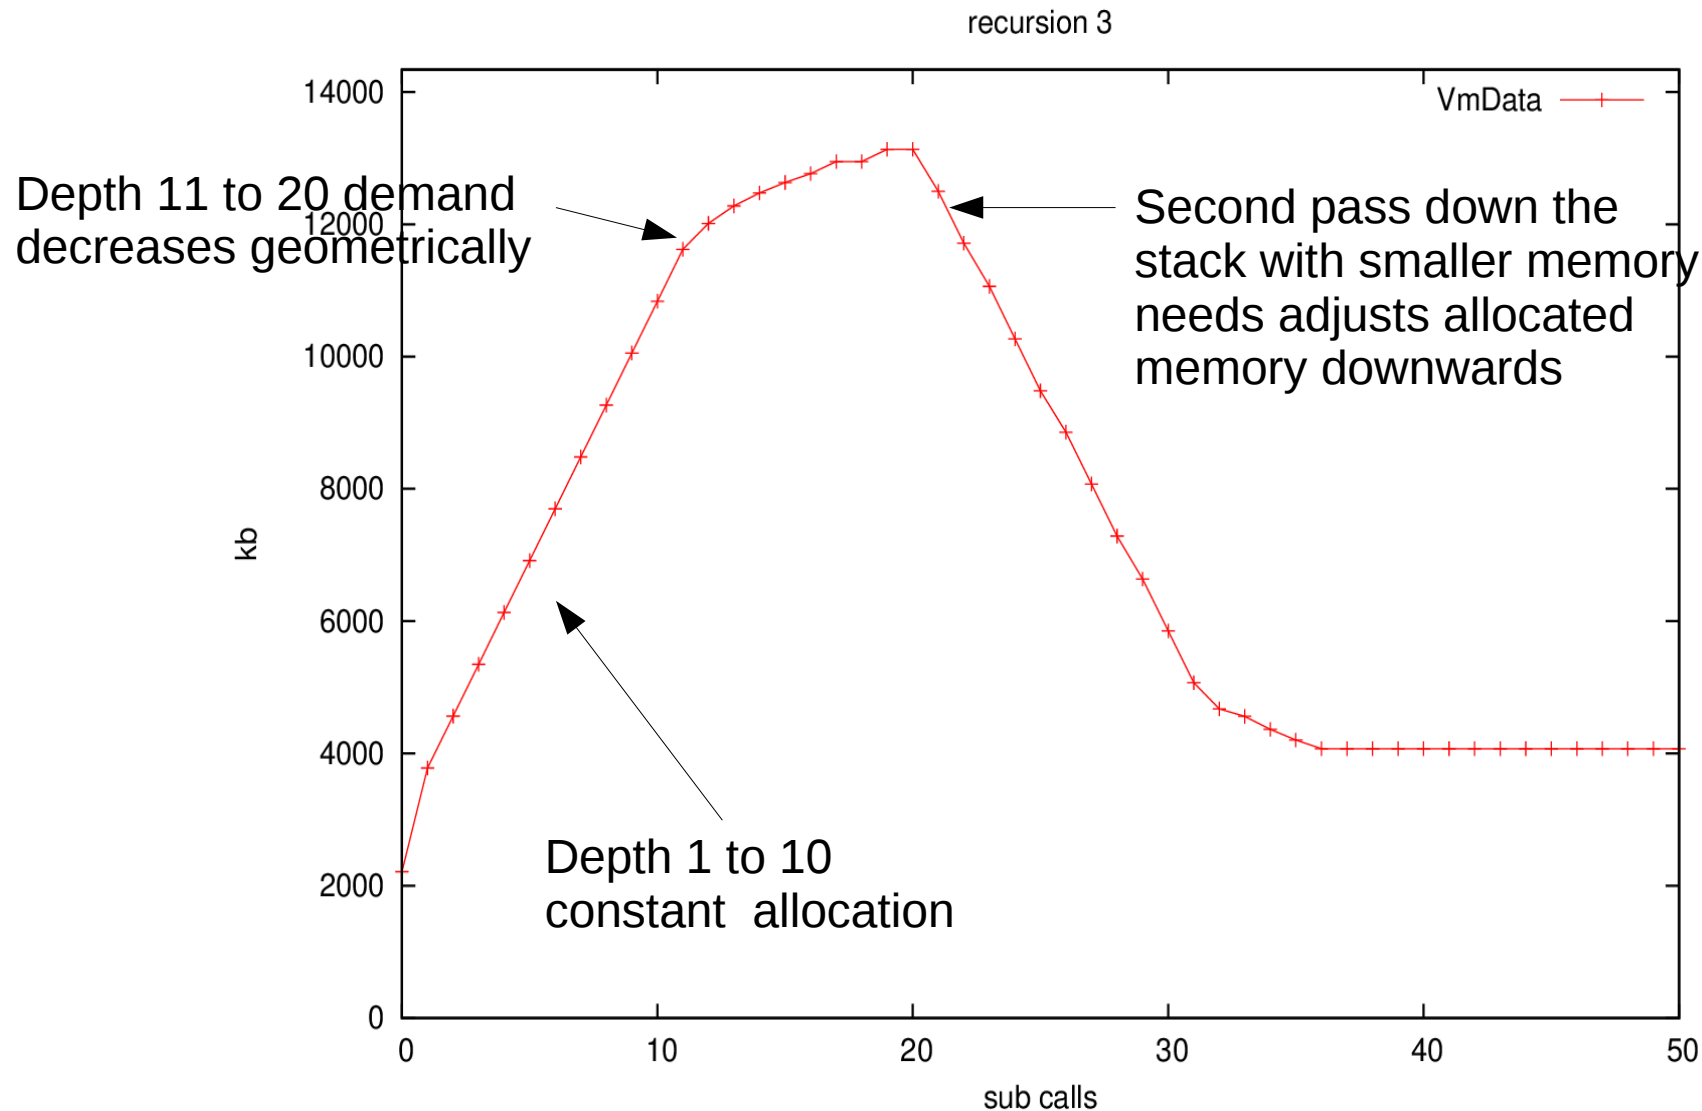
Is memory allocation at given depth based on actual use at that depth?

```
sub get_const {  
    my $descale = shift;  
    $descale -= 10;  
    $descale = 1 if $descale < 1;  
    return chr(ord("a")) x ( 800000 / $descale );  
}
```

```
sub recur {  
    my $cnt = shift;  
    my $limit = shift;  
    return if $cnt >= $limit;  
    $cnt++;  
    ProcVM::print_proc_vm();  
    print "sub $cnt\n";  
    my $s;  
    Dump $s;  
    $s = get_const( $cnt );  
    Dump $s;  
    return recur($cnt, $limit);  
}
```

```
print "cycle 1\n";  
recur(0, 20);  
recur(20, 45);  
recur(45, 50);  
ProcVM::print_proc_vm();
```

Retained allocation based on usage and is adjusted if need decreases next time through

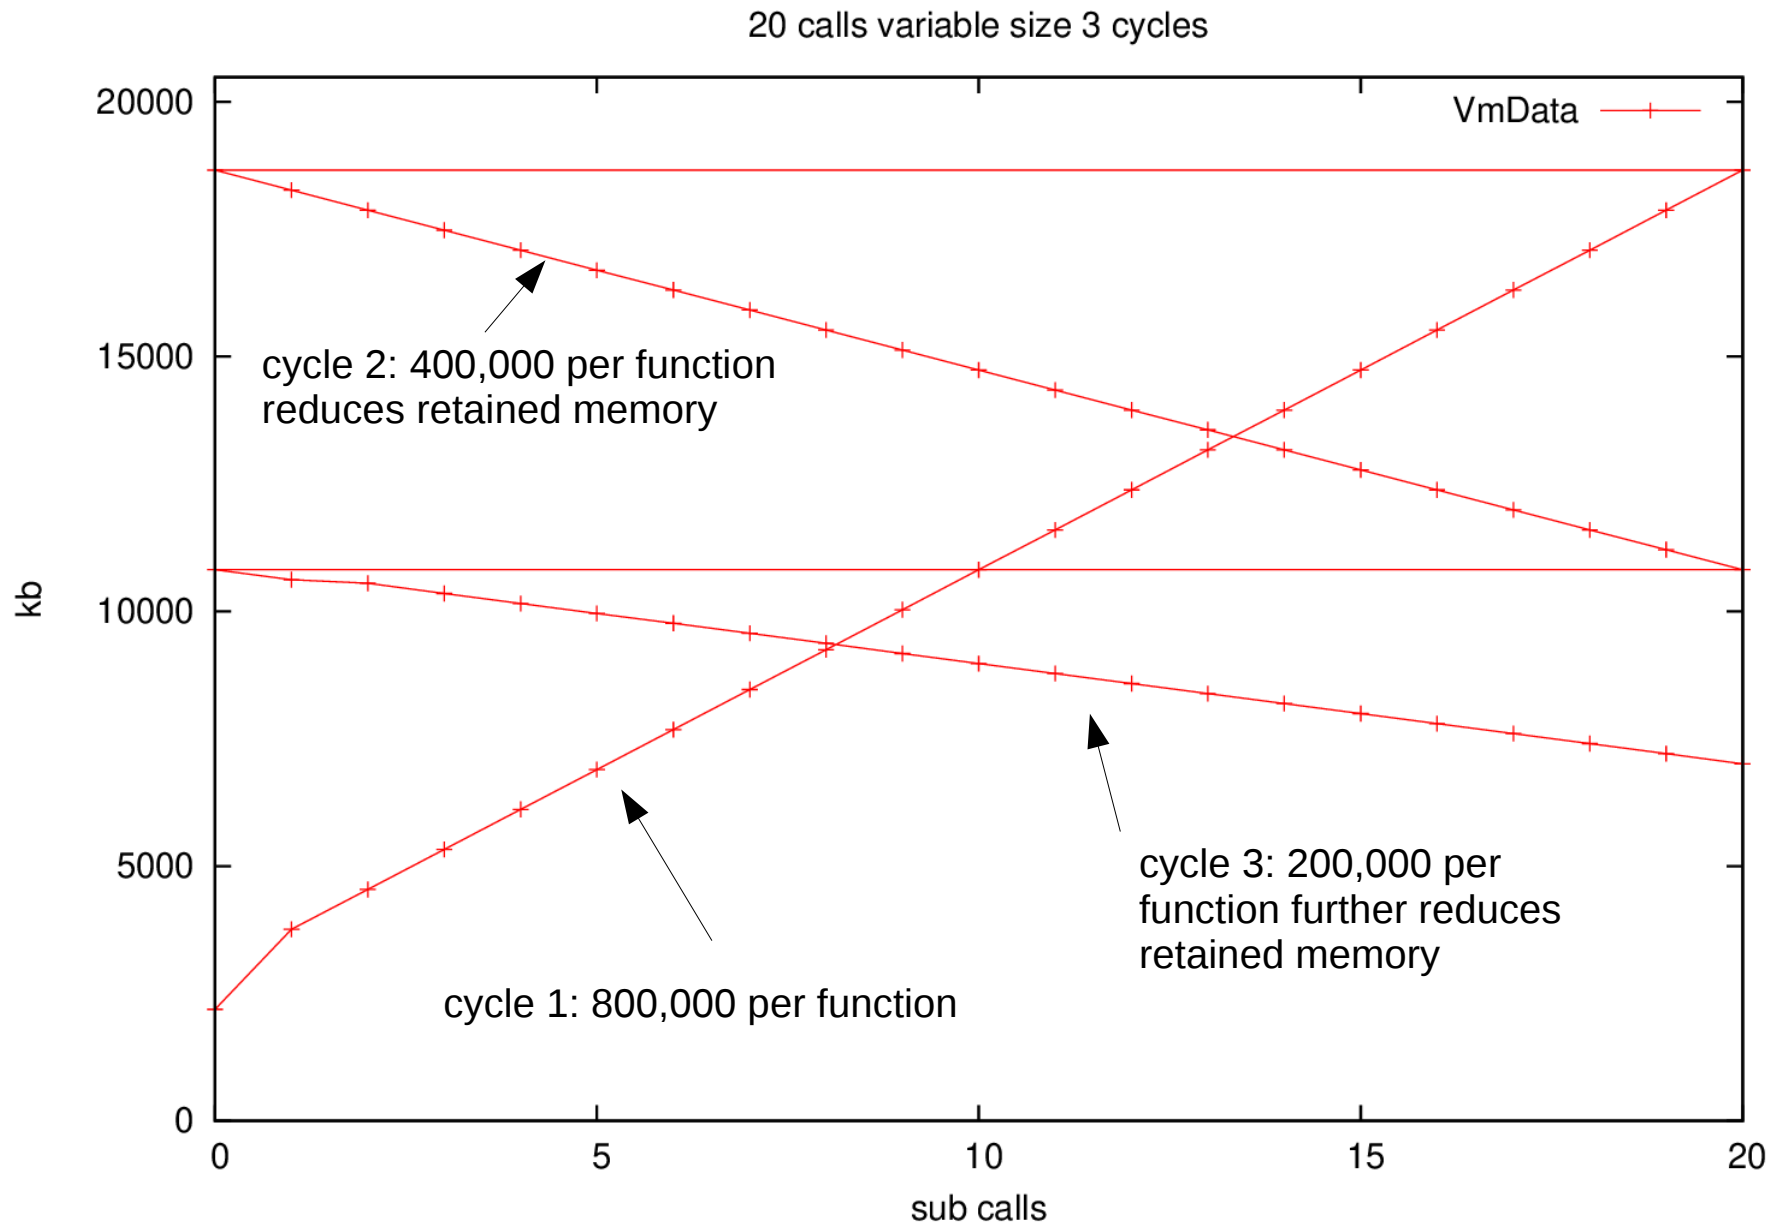


Is memory allocation adjusted in non-recursive case?

Revisit previous test but with variable memory needs

- 20 functions using separate sub to generate big scalar
- cycle 1 full 800,000 character scalar
- cycle 2 half size 400,000 character scalar
- cycle 3 quarter size 200,000 character scalar

Allocated memory adjusts based on usage



What is happening? Is memory retained or not?

```
sub one_big {  
    my $size = shift;  
    return "a" x $size;  
}  
sub big_string {  
    my $size = shift;  
    my $x;  
    Dump $x;  
    $x = one_big($size);  
    Dump $x;  
}
```

```
big_string( 80_000 );  
big_string( 40_000 );
```


In at least some cases where scalar set via return value from another function, retained memory is not reused.

SV = NULL(0x0) at 0x8bdb704

REFCNT = 1

FLAGS = (PADMY)

SV = PV(0x8ba57a8) at 0x8bdb704

REFCNT = 1

FLAGS = (PADMY,POK,pPOK)

PV = 0x8c2a708 "aaaaaaaaaaaaaaaaaa"...\0

CUR = 80000

LEN = 80004

SV = PV(0x8ba57a8) at 0x8bdb704

REFCNT = 1

FLAGS = (PADMY)

PV = 0x8c2a708 "aaaaaaaaaaaaaaaaaa"...\0

CUR = 80000

LEN = 80004

SV = PV(0x8ba57a8) at 0x8bdb704

REFCNT = 1

FLAGS = (PADMY,POK,pPOK)

PV = 0x8c3df90 "aaaaaaaaaaaaaaaaaa"...\0

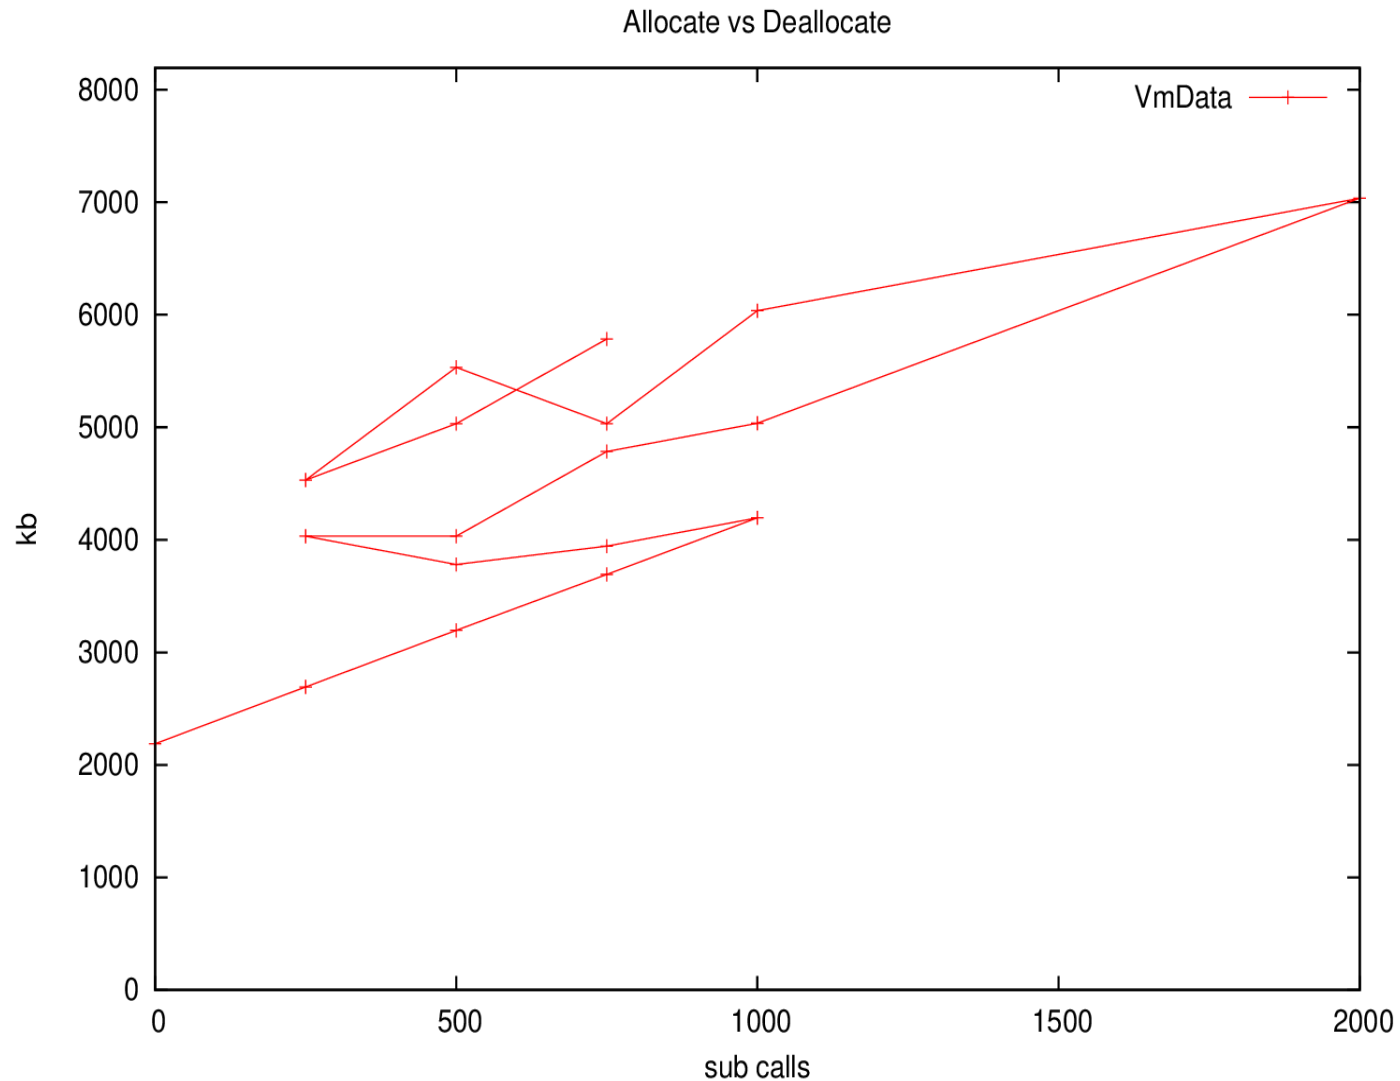
CUR = 40000

LEN = 40004

More complicated example of allocation / deallocation

```
sub big_string {  
    my $size = shift;  
    return "a" x $size;  
}  
  
sub mem_test {  
    my $size = shift;  
    print STDERR "mem_test $size\n";  
    my $s;  
    Dump $s;  
    big_string( $size );  
    $s = big_string( $size );  
    Dump $s;  
    return;  
}  
  
print "cycle 1\n";  
my $i = 1;  
for my $kb (  
    250, 500, 750, 1000,  
    750, 500, 250, 500, 750, 1000, 2000,  
    1000, 750, 500, 250, 500, 750 )  
{  
    ProcVM::print_proc_vm();  
    print "sub $kb\n";  
    mem_test( $kb * 1024 - 4 );  
    $i++  
}  
ProcVM::print_proc_vm();
```

Allocation/deallocation interact with memory management in complicated ways



Perl 5.28.1

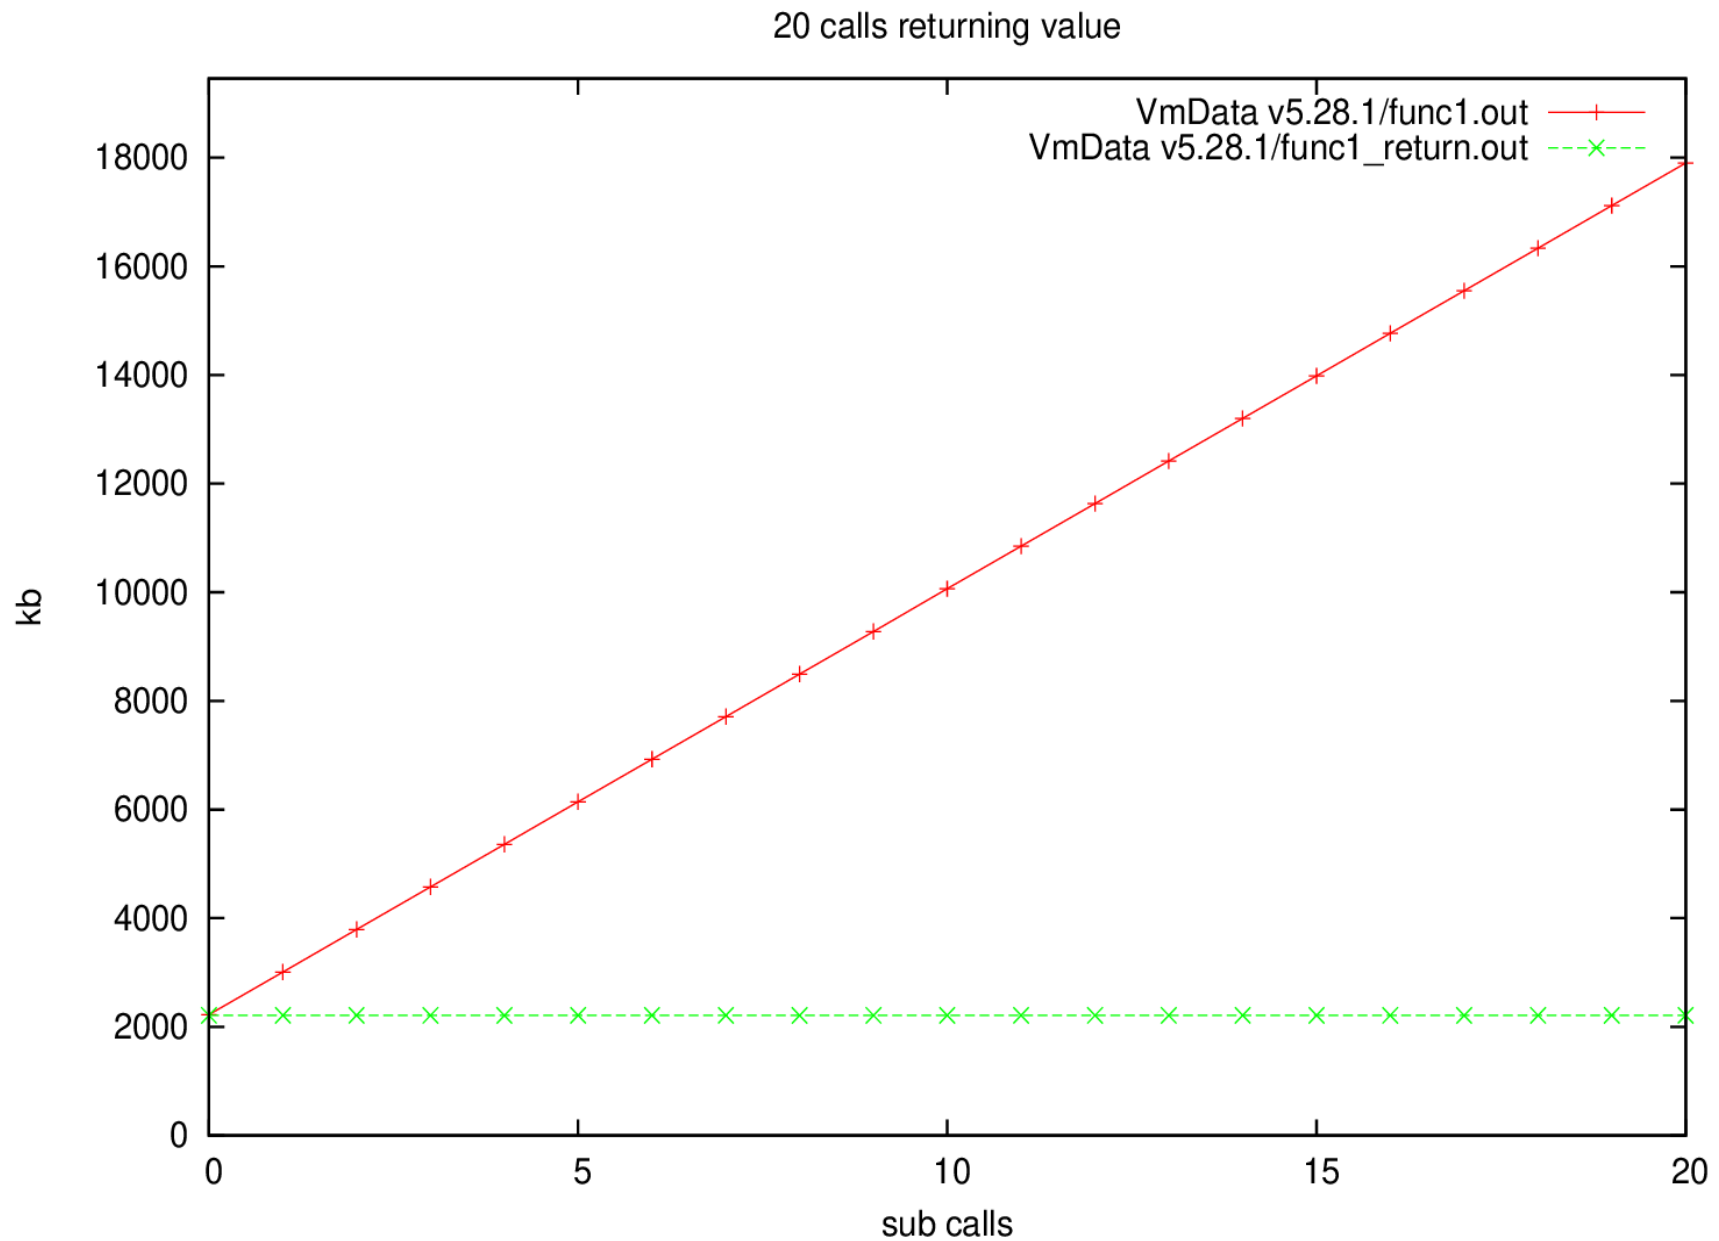
5.28.1 demonstrates many of the same behaviours as 5.18.2 but there are some differences

- Storage allocated to lexical variables still persists between calls
- Since 5.20 perl adds copy on write (COW). Returning a scalar does not immediately create a new copy of a string.

It takes some contortions to reduce persistent memory allocation via returned value

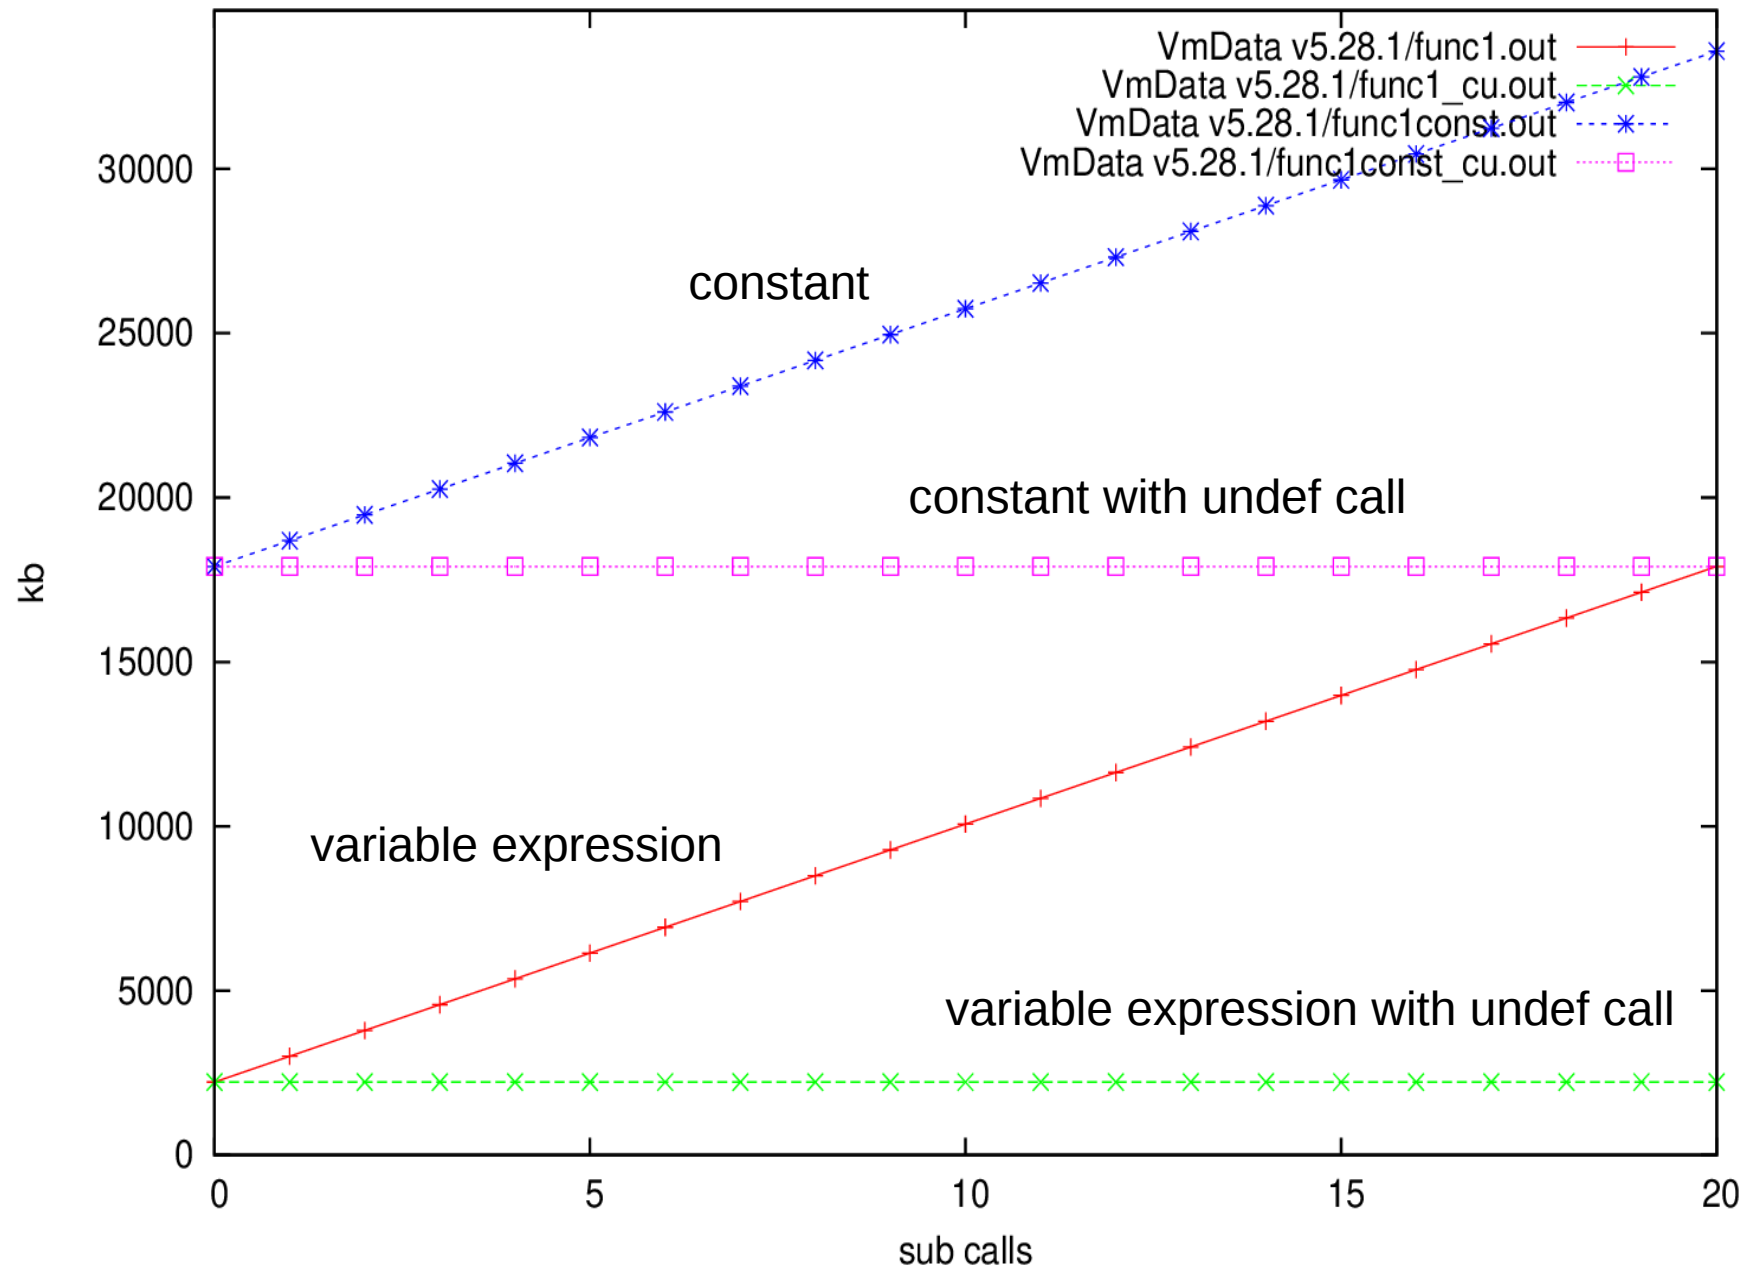
```
sub f1 {  
    my $offset = shift;  
    print "sub 1\n";  
    my $s = (chr(ord("a") + 1 + $offset)) x 800000;  
    return $s;  
}  
...  
  
{  
    my $x = f1(1);  
}  
ProcVM::print_proc_vm();  
{  
    my $x = f2(1);  
}  
ProcVM::print_proc_vm();  
...
```

5.28.1 returning value vs not returning

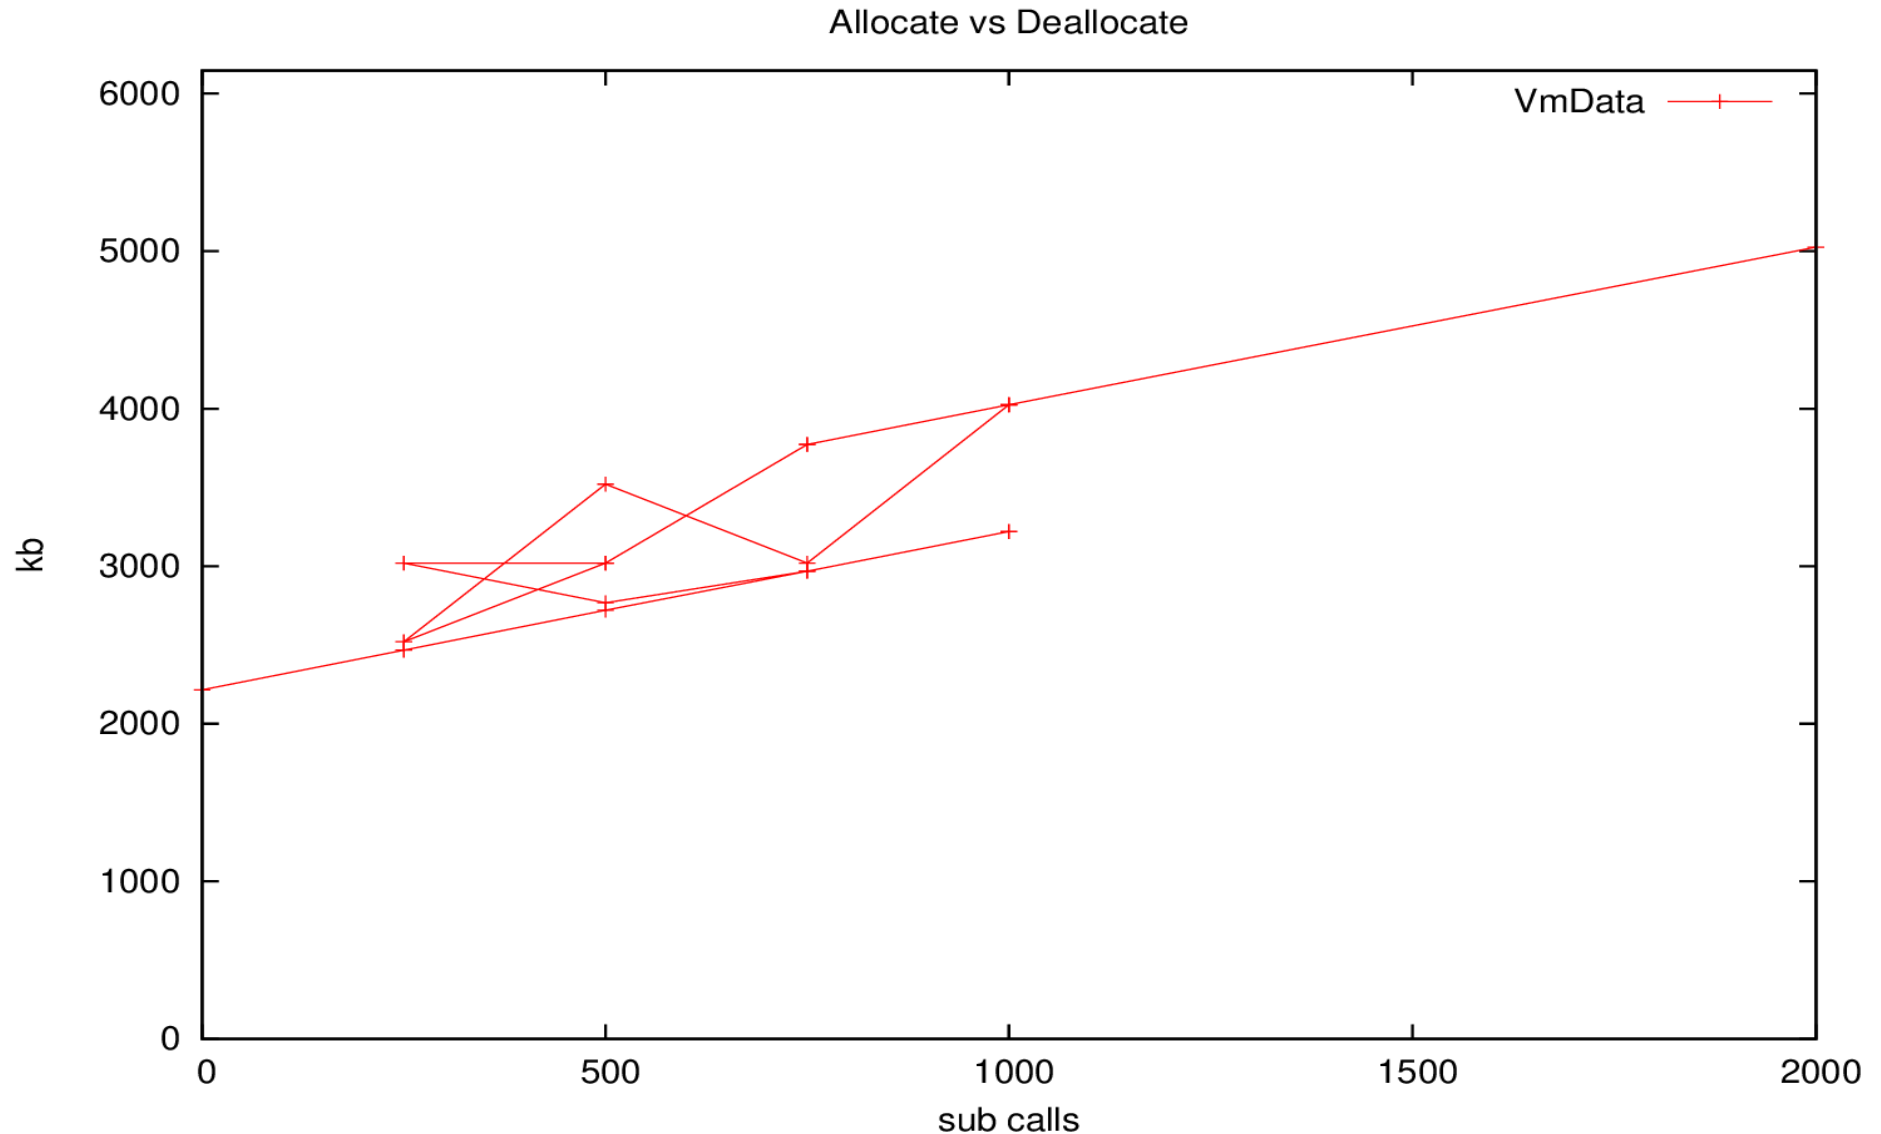


- Perl 5.28 does not consume memory for temporary expressions the way 5.18 did
- but, replacing a variable expression with a constant expression does consume memory for the constant

20 calls vanilla vs nonconst vs const



Allocation/deallocation still interacts with memory management in complicated ways



summary

- Lexical variable memory retained between calls
- Assigning a new value doesn't “usually” release memory.
- Assigning undef doesn't release memory
- Calling undef does release memory
- Hitting a hard memory limit does not force release of memory

- Memory allocated for local variables is released (but we don't want to go there!)
- Return by value may release memory under some circumstances for recent versions of perl
- Returning a reference prevents retention of memory by function
- Memory may also be allocated for evaluating expressions. Less so for recent versions of perl
- Memory for constant expressions may be allocated at “compile” time.

- Closures hand off allocated storage as if they were references. Storage handed off for closures released when ref count hits zero.
- Recursion allocates and retains memory at each level of call depth
- Memory allocation may adjust in subsequent calls, but may depend on specifics of how scalar is populated

Conclusion

Memory management is complicated

*“We should forget about small efficiencies, say about 97% of the time: **premature optimization is the root of all evil**. Yet we should not pass up our opportunities in that critical 3%”*

Donald Knuth

man proc(5)

- VmPeak: Peak virtual memory size.
- VmSize: Virtual memory size.
- VmLck: Locked memory size (see mlock(3)).
- VmHWM: Peak resident set size ("high water mark").
- VmRSS: Resident set size.
- VmData, VmStk, VmExe: Size of data, stack, and text segments.
- VmLib: Shared library code size.
- VmPTE: Page table entries size (since Linux 2.6.10).
- VmSwap: Swapped-out virtual memory size by anonymous private pages; shmem swap usage is not included (since Linux 2.6.34).